



## Practical 1

**Aim: Implement Feed-forward Neural Network and train the network with different optimizers and compare the results.**

### Source Code:

```
# import the necessary packages from
sklearn.preprocessing import LabelBinarizer from
sklearn.metrics import classification_report from
tensorflow.keras.models import Sequential from
tensorflow.keras.layers import Dense from
tensorflow.keras.optimizers import SGD from
tensorflow.keras.datasets import mnist from
tensorflow.keras import backend as K
import matplotlib.pyplot as plt
import numpy as np import
argparse

# construct the argument parse and parse the arguments ap
= argparse.ArgumentParser()
ap.add_argument("-o", "--output", required=True, help="path to the output loss/accuracy plot") args
= vars(ap.parse_args())

# grab the MNIST dataset (if this is your first time using this #
dataset then the 11MB download may take a minute)
print("[INFO] accessing MNIST...")
((trainX, trainY), (testX, testY)) = mnist.load_data()
# each image in the MNIST dataset is represented as a 28x28x1
# image, but in order to apply a standard neural network we must
# first "flatten" the image to be simple list of 28x28=784 pixels
trainX = trainX.reshape((trainX.shape[0], 28 * 28 * 1)) testX =
testX.reshape((testX.shape[0], 28 * 28 * 1))
# scale data to the range of [0, 1] trainX
= trainX.astype("float32") / 255.0
testX = testX.astype("float32") / 255.0

# convert the labels from integers to vectors
lb = LabelBinarizer() trainY =
lb.fit_transform(trainY)
testY = lb.transform(testY)

model = Sequential()
model.add(Dense(256, input_shape=(784,), activation="sigmoid")) model.add(Dense(128,
activation="sigmoid"))
```

```
model.add(Dense(10, activation="softmax"))
```

```
print("[INFO] training network...")
```

```
sgd = SGD(0.01) model.compile(loss="categorical_crossentropy", optimizer=sgd,
metrics=["accuracy"]) H = model.fit(trainX, trainY, validation_data=(testX, testY),
epochs=100, batch_size=128)
```

```
[INFO] training network...
Epoch 1/100
469/469 [=====] - 4s 8ms/step - loss: 2.2798 - accuracy: 0.1753 - val_loss: 2.2416 - val_accuracy: 0.3011
Epoch 2/100
469/469 [=====] - 6s 13ms/step - loss: 2.2074 - accuracy: 0.3917 - val_loss: 2.1643 - val_accuracy: 0.4448
Epoch 3/100
469/469 [=====] - 4s 9ms/step - loss: 2.1176 - accuracy: 0.5159 - val_loss: 2.0556 - val_accuracy: 0.5449
Epoch 4/100
469/469 [=====] - 3s 7ms/step - loss: 1.9872 - accuracy: 0.5869 - val_loss: 1.8971 - val_accuracy: 0.6076
Epoch 5/100
469/469 [=====] - 4s 9ms/step - loss: 1.8047 - accuracy: 0.6278 - val_loss: 1.6870 - val_accuracy: 0.6620
```

```
# evaluate the network
```

```
print("[INFO] evaluating network...") predictions =
```

```
model.predict(testX, batch_size=128)
```

```
print(classification_report(testY.argmax(axis=1), predictions.argmax(axis=1),
target_names=[str(x) for x in lb.classes_]))
```

```
plt.style.use("ggplot")
```

```
plt.figure() plt.plot(np.arange(0, 100), H.history["loss"],
label="train_loss") plt.plot(np.arange(0, 100), H.history["val_loss"],
label="val_loss") plt.plot(np.arange(0, 100), H.history["accuracy"],
label="train_acc") plt.plot(np.arange(0, 100),
H.history["val_accuracy"], label="val_acc") plt.title("Training Loss
and Accuracy") plt.xlabel("Epoch #") plt.ylabel("Loss/Accuracy")
plt.legend()
```



```
# import the necessary packages
```

```

from sklearn.preprocessing import LabelBinarizer
from sklearn.metrics import classification_report
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.datasets import cifar10
import matplotlib.pyplot as plt
import numpy as np
import argparse

print("[INFO] loading CIFAR-10 data...")
((trainX, trainY), (testX, testY)) = cifar10.load_data()
trainX = trainX.astype("float") / 255.0
testX = testX.astype("float") / 255.0
trainX = trainX.reshape((trainX.shape[0], 3072))
testX = testX.reshape((testX.shape[0], 3072))

lb = LabelBinarizer()
trainY = lb.fit_transform(trainY)
testY = lb.transform(testY)

# initialize the label names for the CIFAR-10 dataset
labelNames = ["airplane", "automobile", "bird", "cat", "deer", "dog", "frog", "horse", "ship", "truck"]

model = Sequential()
model.add(Dense(1024, input_shape=(3072,), activation="relu"))
model.add(Dense(512, activation="relu"))
model.add(Dense(10, activation="softmax"))

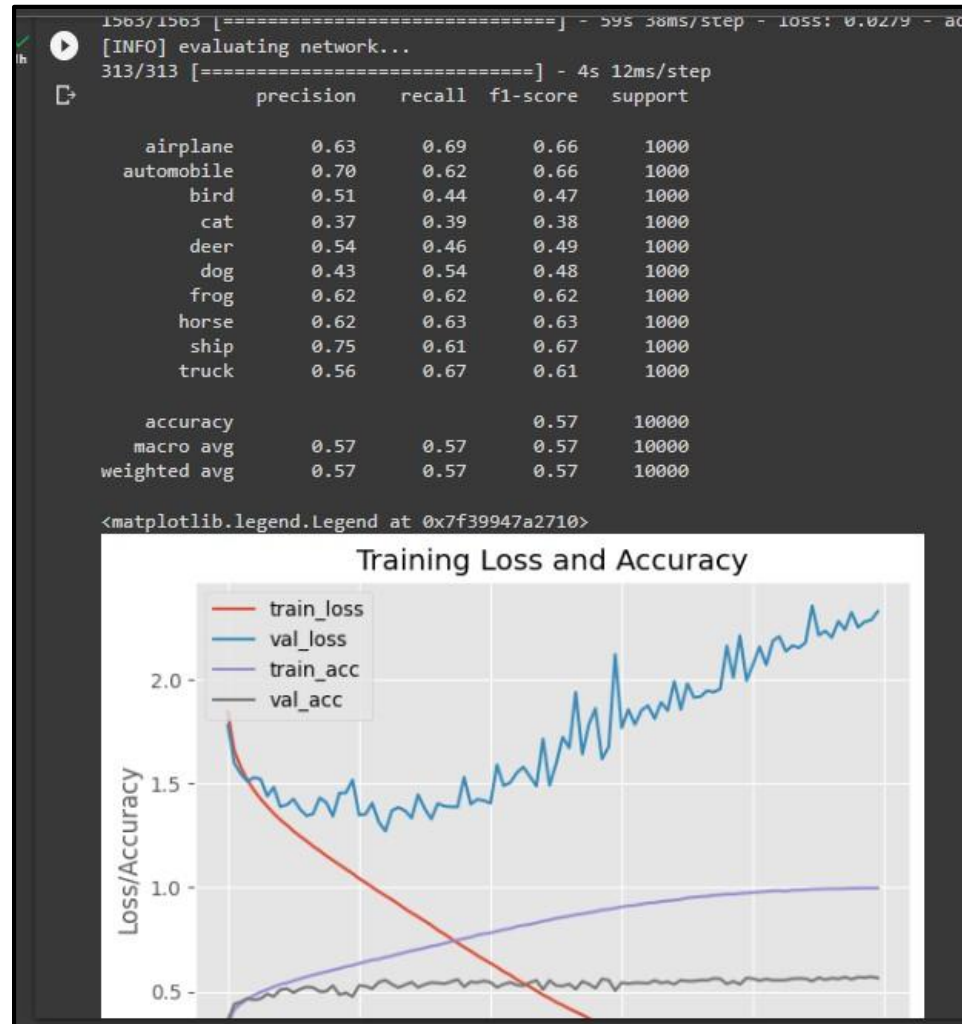
print("[INFO] training network...")
sgd = SGD(0.01)
model.compile(loss="categorical_crossentropy", optimizer=sgd, metrics=["accuracy"])
H = model.fit(trainX, trainY, validation_data=(testX, testY), epochs=100, batch_size=32)

print("[INFO] evaluating network...")
predictions = model.predict(testX, batch_size=32)
print(classification_report(testY.argmax(axis=1), predictions.argmax(axis=1), target_names=labelNames))

plt.style.use("ggplot")
plt.figure()
plt.plot(np.arange(0, 100), H.history["loss"], label="train_loss")
plt.plot(np.arange(0, 100), H.history["val_loss"], label="val_loss")
plt.plot(np.arange(0, 100), H.history["accuracy"], label="train_acc")
plt.plot(np.arange(0, 100), H.history["val_accuracy"], label="val_acc")

```

```
plt.title("Training Loss and Accuracy") plt.xlabel("Epoch #")
plt.ylabel("Loss/Accuracy") plt.legend()
```



```
from sklearn.preprocessing import LabelBinarizer
from sklearn.metrics import classification_report
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.datasets import cifar10
import matplotlib.pyplot as plt
import numpy as np
import argparse
```

```
# load the training and testing data, scale it into the range [0, 1],
# then reshape the design matrix print("[INFO]
loading CIFAR-10 data...") ((trainX, trainY), (testX,
testY)) = cifar10.load_data() trainX =
trainX.astype("float") / 255.0 testX =
testX.astype("float") / 255.0
```

```

trainX = trainX.reshape((trainX.shape[0], 3072)) testX
= testX.reshape((testX.shape[0], 3072))

# convert the labels from integers to vectors
lb = LabelBinarizer() trainY =
lb.fit_transform(trainY) testY =
lb.transform(testY)
# initialize the label names for the CIFAR-10 dataset labelNames
= ["airplane", "automobile", "bird", "cat", "deer",
    "dog", "frog", "horse", "ship", "truck"]

# define the 3072-1024-512-10 architecture using Keras model
= Sequential()
model.add(Dense(1024, input_shape=(3072,), activation="relu")) model.add(Dense(512,
activation="relu"))
model.add(Dense(10, activation="softmax"))

# train the model using SGD print("[INFO]
training network...")
sgd = SGD(0.01) model.compile(loss="categorical_crossentropy", optimizer=sgd,
metrics=["accuracy"]) H = model.fit(trainX, trainY, validation_data=(testX, testY),
epochs=100, batch_size=32)

print("[INFO] evaluating network...") predictions
= model.predict(testX, batch_size=32)
print(classification_report(testY.argmax(axis=1),
    predictions.argmax(axis=1), target_names=labelNames))

# plot the training loss and accuracy plt.style.use("ggplot")
plt.figure()
plt.plot(np.arange(0, 100), H.history["loss"], label="train_loss")
plt.plot(np.arange(0, 100), H.history["val_loss"], label="val_loss")
plt.plot(np.arange(0, 100), H.history["accuracy"], label="train_acc")
plt.plot(np.arange(0, 100), H.history["val_accuracy"], label="val_acc")
plt.title("Training Loss and Accuracy") plt.xlabel("Epoch #")
plt.ylabel("Loss/Accuracy") plt.legend()
plt.savefig(args["output"])

```

## Output:

```
Epoch 99/100
1563/1563 [=====] - 58s 37ms/step - loss: 0.0304 -
Epoch 100/100
1563/1563 [=====] - 59s 38ms/step - loss: 0.0256 -
[INFO] evaluating network...
313/313 [=====] - 5s 15ms/step
```

	precision	recall	f1-score	support
airplane	0.58	0.70	0.63	1000
automobile	0.69	0.69	0.69	1000
bird	0.49	0.38	0.43	1000
cat	0.39	0.38	0.39	1000
deer	0.45	0.50	0.47	1000
dog	0.47	0.44	0.45	1000
frog	0.60	0.63	0.62	1000
horse	0.64	0.60	0.62	1000
ship	0.69	0.71	0.70	1000
truck	0.61	0.60	0.61	1000
accuracy			0.56	10000
macro avg	0.56	0.56	0.56	10000
weighted avg	0.56	0.56	0.56	10000



## Practical 2

**Aim: Write a Program to implement regularization to prevent the model from overfitting**

**Source Code:**

```
import numpy as np
import pandas as pd
from sklearn import metrics
from sklearn.linear_model import Lasso

df_train = pd.read_csv('/content/sample_data/train.csv')
df_test = pd.read_csv('/content/sample_data/test.csv')

df_train = df_train.dropna()
df_test = df_test.dropna()

x_train = df_train['x']
y_train = df_train['y']
x_train = x_train.values.reshape(-1,1)
y_train = y_train.values.reshape(-1,1)

x_test = df_test['x']
y_test = df_test['y']
x_test = x_test.values.reshape(-1,1)
y_test = y_test.values.reshape(-1,1)

lasso = Lasso()

lasso.fit(x_train, y_train)

print("Lasso Train RMSE:", np.round(np.sqrt(metrics.mean_squared_error(y_train,
lasso.predict(x_train))), 5))
print("Lasso Test RMSE:", np.round(np.sqrt(metrics.mean_squared_error(y_test, lasso.predict(x_test))),
5))

import numpy as np
import pandas as pd
from sklearn import metrics
from sklearn.linear_model import Ridge

df_train = pd.read_csv('/content/sample_data/train.csv')
df_test = pd.read_csv('/content/sample_data/test.csv')
```



```
df_train = df_train.dropna()
df_test = df_test.dropna()

x_train = df_train['x']
x_train = x_train.values.reshape(-1,1) y_train
= df_train['y']
y_train = y_train.values.reshape(-1,1)

x_test = df_test['x']
x_test = x_test.values.reshape(-1,1) y_test
= df_test['y']
y_test = y_test.values.reshape(-1,1)

ridge = Ridge()

ridge.fit(x_train, y_train)
print("Ridge Train RMSE:", np.round(np.sqrt(metrics.mean_squared_error(y_train,
ridge.predict(x_train))), 5))
print("Ridge Test RMSE:", np.round(np.sqrt(metrics.mean_squared_error(y_test, ridge.predict(x_test))),
5))
```

### Output:

```
ridge.fit(x_train, y_train)
print("Ridge Train RMSE:", np.round(np.sqrt(metrics.mean_squared_error(y_train, ridge.predict(x_train))), 5))
print("Ridge Test RMSE:", np.round(np.sqrt(metrics.mean_squared_error(y_test, ridge.predict(x_test))), 5))

Lasso Train RMSE: 2.80516
Lasso Test RMSE: 3.07592
Ridge Train RMSE: 2.80495
Ridge Test RMSE: 3.07131
```

## Practical 3

**Aim: Implement deep learning for recognizing classes for datasets like CIFAR-10 images for previously unseen images and assign them to one of the 10 classes**

**Source Code:**

```
import tensorflow as tf
from tensorflow.keras import datasets, layers, models
import matplotlib.pyplot as plt import numpy as np

(X_train, y_train), (X_test,y_test) = datasets.cifar10.load_data()
X_train.shape

X_test.shape

y_train.shape

y_train[:5]

y_train = y_train.reshape(-1,) y_train[:5]

y_test = y_test.reshape(-1,)

classes = ["airplane","automobile","bird","cat","deer","dog","frog","horse","ship","truck"]

def plot_sample(X, y, index):
    plt.figure(figsize = (15,2))
    plt.imshow(X[index])
    plt.xlabel(classes[y[index]])

plot_sample(X_train, y_train, 0)

X_train = X_train / 255.0
X_test = X_test / 255.0

ann = models.Sequential([
    layers.Flatten(input_shape=(32,32,3)),
    layers.Dense(3000, activation='relu'), layers.Dense(1000,
activation='relu'), layers.Dense(10, activation='softmax')
])
```

```

ann.compile(optimizer='SGD',
            loss='sparse_categorical_crossentropy',
            metrics=['accuracy'])

ann.fit(X_train, y_train, epochs=5)

from sklearn.metrics import confusion_matrix , classification_report
import numpy as np
y_pred = ann.predict(X_test)
y_pred_classes = [np.argmax(element) for element in y_pred]

print("Classification Report: \n", classification_report(y_test, y_pred_classes))

cnn = models.Sequential([
    layers.Conv2D(filters=32, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),

    layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])

cnn.compile(optimizer='adam',
            loss='sparse_categorical_crossentropy',
            metrics=['accuracy'])

cnn.fit(X_train, y_train, epochs=10)

cnn.evaluate(X_test,y_test)

y_pred = cnn.predict(X_test)
y_pred[:5]

y_classes = [np.argmax(element) for element in y_pred]
y_classes[:5]

y_test[:5]

plot_sample(X_test, y_test,3)

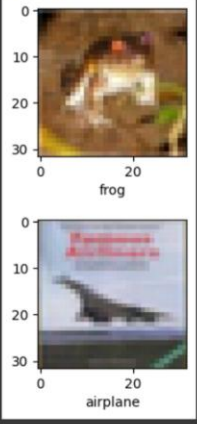
```

```
classes[y_classes[3]]
```

```
classes[y_classes[3]]
```

## Output:

```
1563/1563 [=====] - 69s 44ms/step - loss: 0.8212 - accuracy: 0.7143
Epoch 7/10
1563/1563 [=====] - 69s 44ms/step - loss: 0.7759 - accuracy: 0.7290
Epoch 8/10
1563/1563 [=====] - 70s 45ms/step - loss: 0.7329 - accuracy: 0.7456
Epoch 9/10
1563/1563 [=====] - 69s 44ms/step - loss: 0.6932 - accuracy: 0.7589
Epoch 10/10
1563/1563 [=====] - 70s 44ms/step - loss: 0.6592 - accuracy: 0.7697
313/313 [=====] - 7s 20ms/step - loss: 0.9105 - accuracy: 0.6990
313/313 [=====] - 4s 12ms/step
'airplane'
```



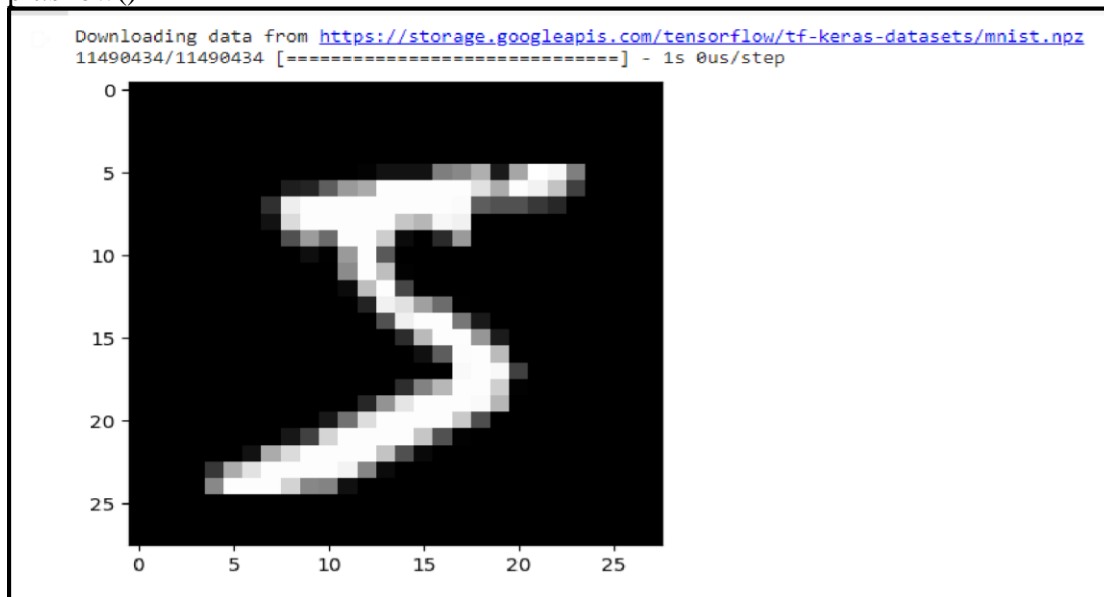
## Practical 4

**Aim: Implement deep learning for the Prediction of the autoencoder from the test data (e.g. MNIST data set)**

**Source Code:**

```
!pip install tensorflow-gpu==2.0.0b1 from
tensorflow.keras.datasets import mnist from
tensorflow.keras.layers import Dense, Input, Flatten,\
    Reshape, LeakyReLU as LR,\
    Activation, Dropout
from tensorflow.keras.models import Model, Sequential from
matplotlib import pyplot as plt
from IPython import display # If using IPython, Colab or Jupyter import
numpy as np
```

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train = x_train/255.0 x_test = x_test/255.0
# Plot image data from x_train plt.imshow(x_train[0],
cmap = "gray")
plt.show()
```



```
LATENT_SIZE = 32
encoder = Sequential([
    Flatten(input_shape = (28, 28)),
    Dense(512),
    LR(),
```

```

        Dropout(0.5),
        Dense(256),
        LR(),
        Dropout(0.5),
        Dense(128),
        LR(),
        Dropout(0.5),
        Dense(64),
        LR(),
        Dropout(0.5),
        Dense(LATENT_SIZE, activation="sigmoid"),
    ])

```

```

decoder = Sequential([
    Dense(64, input_shape = (LATENT_SIZE,)),
    LR(),
    Dropout(0.5),
    Dense(128),
    LR(),
    Dropout(0.5),
    Dense(256),
    LR(),
    Dropout(0.5),
    Dense(512),
    LR(),
    Dropout(0.5),
    Dense(784),
    Activation("sigmoid"),
    Reshape((28, 28))
])

```

```

img = Input(shape = (28, 28)) latent_vector =
encoder(img) output = decoder(latent_vector) model =
Model(inputs = img, outputs = output)
model.compile("nadam", loss = "binary_crossentropy")
EPOCHS = 100
#Only do plotting if you have IPython, Jupyter, or using Colab for
epoch in range(EPOCHS):
    fig, axs = plt.subplots(4, 4, figsize=(4,4))
    rand = x_test[np.random.randint(0, 10000, 16)].reshape((4, 4, 1, 28, 28))


```

```
display.clear_output() # If you imported display from IPython
```

```
for i in range(4):  
for j in range(4):  
    axs[i, j].imshow(model.predict(rand[i, j])[0], cmap = "gray")  
axs[i, j].axis("off")
```

```
plt.subplots_adjust(wspace = 0, hspace = 0)  
plt.show()  
print("-----", "EPOCH", epoch, "-----")  
model.fit(x_train, x_train, batch_size = 64)
```

```
1/1 [=====] - 0s 27ms/step  
1/1 [=====] - 0s 33ms/step  
1/1 [=====] - 0s 30ms/step  
1/1 [=====] - 0s 34ms/step  
1/1 [=====] - 0s 30ms/step  
1/1 [=====] - 0s 28ms/step  
1/1 [=====] - 0s 20ms/step  
1/1 [=====] - 0s 20ms/step  
1/1 [=====] - 0s 19ms/step  
1/1 [=====] - 0s 18ms/step  
1/1 [=====] - 0s 21ms/step  
1/1 [=====] - 0s 17ms/step  
1/1 [=====] - 0s 18ms/step  
1/1 [=====] - 0s 18ms/step  
1/1 [=====] - 0s 21ms/step  
1/1 [=====] - 0s 18ms/step
```



```
----- EPOCH 99 -----  
938/938 [=====] - 6s 7ms/step - loss: 0.1837
```

## Practical 5

**Aim: Implement Convolutional Neural Network for Digit Recognition on the MNIST Dataset.**

### Source Code:

```
# baseline cnn model for mnist from numpy import
mean from numpy import std from matplotlib
import pyplot as plt from sklearn.model_selection
import KFold from tensorflow.keras.datasets import
mnist from tensorflow.keras.utils import
to_categorical from tensorflow.keras.models import
Sequential from tensorflow.keras.layers import
Conv2D from tensorflow.keras.layers import
MaxPooling2D from tensorflow.keras.layers import
Dense from tensorflow.keras.layers import Flatten
from tensorflow.keras.optimizers import SGD

# load train and test dataset def
load_dataset():
    # load dataset
    (trainX, trainY), (testX, testY) = mnist.load_data() #
    reshape dataset to have a single channel trainX =
    trainX.reshape((trainX.shape[0], 28, 28, 1)) testX =
    testX.reshape((testX.shape[0], 28, 28, 1))
    # one hot encode target values
    trainY = to_categorical(trainY) testY
    = to_categorical(testY)
    return trainX, trainY, testX, testY

# scale pixels def prep_pixels(train,
test): # convert from integers to
floats train_norm =
train.astype('float32') test_norm =
test.astype('float32') # normalize to
range 0-1 train_norm = train_norm
/ 255.0 test_norm = test_norm /
255.0
# return normalized images
return train_norm, test_norm
```



```

# define cnn model def
def define_model():
    model = Sequential()
    model.add(Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform',
input_shape=(28, 28, 1))) model.add(MaxPooling2D((2, 2)))
    model.add(Flatten())
    model.add(Dense(100, activation='relu', kernel_initializer='he_uniform'))
    model.add(Dense(10, activation='softmax'))
    # compile model
    opt = SGD(learning_rate=0.01, momentum=0.9)
    model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
    return model

# evaluate a model using k-fold cross-validation
def evaluate_model(dataX, dataY, n_folds=5):
    scores, histories = list(), list() # prepare cross
    validation
    kfold = KFold(n_folds, shuffle=True, random_state=1)
    # enumerate splits for train_ix, test_ix in
    kfold.split(dataX):
        # define model model =
        define_model()
        # select rows for train and test
        trainX, trainY, testX, testY = dataX[train_ix], dataY[train_ix], dataX[test_ix], dataY[test_ix]
        # fit model
        history = model.fit(trainX, trainY, epochs=10, batch_size=32, validation_data=(testX, testY),
        verbose=0) # evaluate model
        _, acc = model.evaluate(testX, testY, verbose=0)
        print('> %.3f' % (acc * 100.0))
        # stores scores
        scores.append(acc)
        histories.append(history)
    return scores, histories

def summarize_diagnostics(histories):
    for i in range(len(histories)):
        # plot loss plt.subplot(2, 1, 1) plt.title('Cross Entropy Loss')
        plt.plot(histories[i].history['loss'], color='blue', label='train')
        plt.plot(histories[i].history['val_loss'], color='orange', label='test')
        # plot accuracy
        plt.subplot(2, 1, 2)

```

```

plt.title('Classification Accuracy')
plt.plot(histories[i].history['accuracy'], color='blue', label='train')
plt.plot(histories[i].history['val_accuracy'], color='orange', label='test')
plt.show()

# summarize model performance def
summarize_performance(scores):
    # print summary
    print('Accuracy: mean=%.3f std=%.3f, n=%d' % (mean(scores)*100, std(scores)*100, len(scores)))
    # box and whisker plots of results
    plt.boxplot(scores)
    plt.show()

def run_test_harness():
    # load dataset
    trainX, trainY, testX, testY = load_dataset()
    # prepare pixel data
    trainX, testX = prep_pixels(trainX, testX)
    # evaluate model
    scores, histories = evaluate_model(trainX, trainY)
    # learning curves
    summarize_diagnostics(histories) #
    summarize_estimated_performance
    summarize_performance(scores)

run_test_harness()

```

### Output:

```

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [=====] - 0s 0us/step
> 98.442
> 98.733

```

## Practical 6

**Aim: Write a program to implement Transfer Learning on the suitable dataset (e.g. classify the cats versus dogs dataset from Kaggle).**

### Source Code:

```
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import os
import zipfile

from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.applications import VGG16

filename = "cats_and_dogs_filtered.zip"

with zipfile.ZipFile("/content/drive/MyDrive/cats_and_dogs_filtered.zip", "r") as zip_ref :
    zip_ref.extractall()

train_dir = os.path.join(os.getcwd(),"cats_and_dogs_filtered","train")
validation_dir = os.path.join(os.getcwd(),"cats_and_dogs_filtered","validation")

train_datagen = ImageDataGenerator(rescale=1./255,rotation_range=20,width_shift_range=0.2,height_shift_range=0.2,shear_range=0.2,zoom_range=0.2,horizontal_flip=True)

validate_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(train_dir,target_size=(150,150),batch_size=20,class_mode="binary")

validation_generator = validate_datagen.flow_from_directory(validation_dir,target_size=(150,150),batch_size=20,class_mode="binary")

conv_base = VGG16(weights="imagenet",include_top=False, input_shape=(150,150,3))

conv_base.trainable = False

model = tf.keras.models.Sequential()
model.add(conv_base)
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(256,activation = "relu"))
model.add(tf.keras.layers.Dropout(0.5))
model.add(tf.keras.layers.Dense(1,activation = "sigmoid"))
```

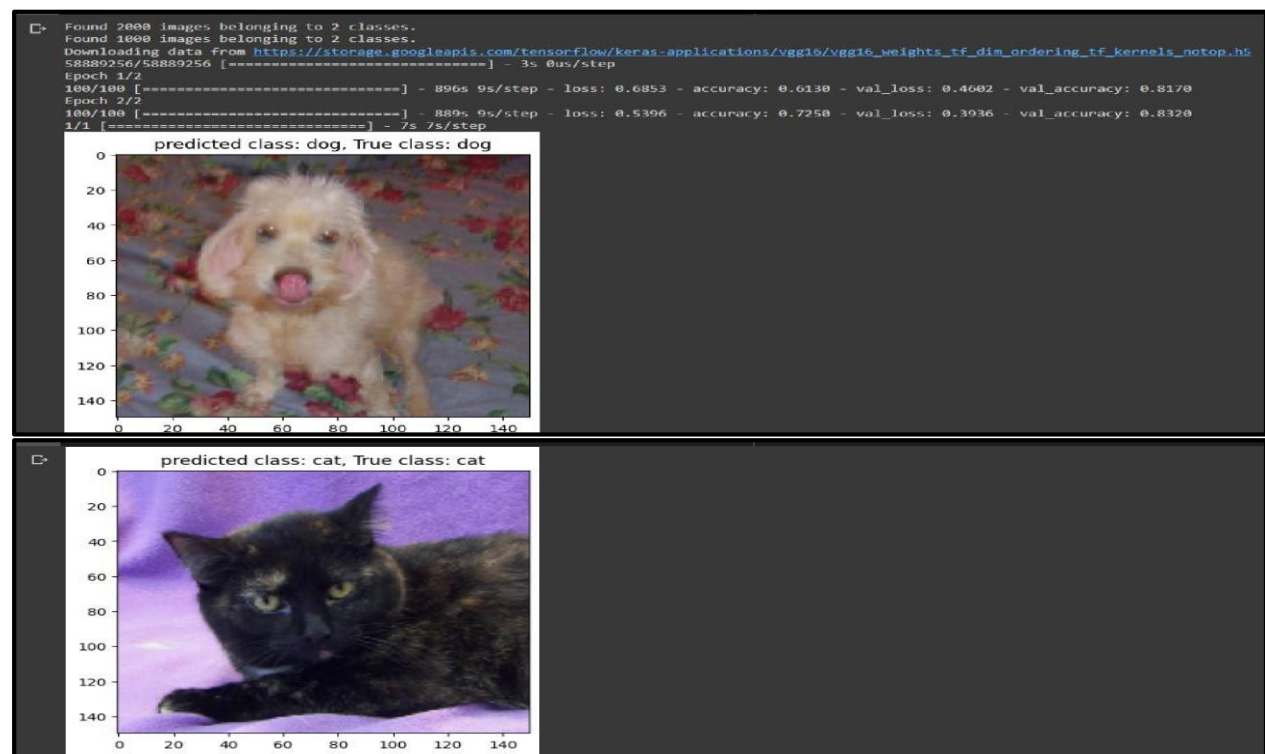
```
model.compile(loss="binary_crossentropy",optimizer=tf.keras.optimizers.RMSprop(learning_rate=2e-5),metrics=["accuracy"])
```

```
history = model.fit(train_generator,steps_per_epoch=100, epochs=2,
validation_data=validation_generator,validation_steps=50)
```

```
x, y_true = next(validation_generator) y_pred = model.predict(x) class_names
= ['cat', 'dog'] for i in range(len(x)): plt.imshow(x[i]) plt.title(f'predicted
class: {class_names[int(round(y_pred[i][0]))]}, True class:
{class_names[int(y_true[i])]}')
plt.show()
```


```
acc = history.history["accuracy"] val_acc
= history.history["val_accuracy"] loss =
history.history["loss"] val_loss =
history.history["val_loss"]
```

```
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, "bo", label="Training acc")
plt.plot(epochs, val_acc, "b", label="validation acc")
plt.title("Training and validation accuracy")
plt.legend() plt.figure() plt.plot(epochs, loss, "bo",
label="Training loss") plt.plot(epochs, val_loss, "b",
label="validation loss") plt.title("Training and
validation loss") plt.legend() plt.show() Output:
```






Training and validation accuracy



Training and validation loss



## Practical 7

**Aim: Write a program for the Implementation of a Generative Adversarial Network for generating synthetic shapes (like digits)**

**Source Code:** from numpy import  
expand\_dims from numpy import ones  
from numpy import zeros from  
numpy.random import rand from  
numpy.random import randint from  
keras.datasets.mnist import load\_data from  
keras.optimizers import Adam from  
keras.models import Sequential from  
keras.layers import Dense from  
keras.layers import Conv2D from  
keras.layers import Flatten from  
keras.layers import Dropout  
from keras.layers import LeakyReLU

```
# define the standalone discriminator model def
define_discriminator(in_shape=(28,28,1)):
model = Sequential()
model.add(Conv2D(64, (3,3), strides=(2, 2), padding='same', input_shape=in_shape))
model.add(LeakyReLU(alpha=0.2)) model.add(Dropout(0.4))
model.add(Conv2D(64, (3,3), strides=(2, 2), padding='same'))
model.add(LeakyReLU(alpha=0.2))
model.add(Dropout(0.4)) model.add(Flatten())
model.add(Dense(1, activation='sigmoid'))
# compile model
opt = Adam(lr=0.0002, beta_1=0.5)
model.compile(loss='binary_crossentropy', optimizer=opt, metrics=['accuracy'])
return model
```

```
# load and prepare mnist training images def
load_real_samples():
# load mnist dataset
(trainX, _), (_, _) = load_data()
# expand to 3d, e.g. add channels dimension
X = expand_dims(trainX, axis=-1)
# convert from unsigned ints to floats
X = X.astype('float32')
```

```

# scale from [0,255] to [0,1]
X = X / 255.0 return X

# select real samples def
generate_real_samples(dataset, n_samples):
# choose random instances
ix = randint(0, dataset.shape[0], n_samples)
# retrieve selected images
X = dataset[ix]
# generate 'real' class labels (1)
y = ones((n_samples, 1))
return X, y

# generate n fake samples with class labels
def generate_fake_samples(n_samples): #
generate uniform random numbers in [0,1]
X = rand(28 * 28 * n_samples)
# reshape into a batch of grayscale images
X = X.reshape((n_samples, 28, 28, 1))
# generate 'fake' class labels (0) y =
zeros((n_samples, 1))
return X, y

# train the discriminator model def train_discriminator(model,
dataset, n_iter=100, n_batch=256):
half_batch = int(n_batch / 2)
# manually enumerate epochs
for i in range(n_iter):
# get randomly selected 'real' samples
X_real, y_real = generate_real_samples(dataset, half_batch)
# update discriminator on real samples
_, real_acc = model.train_on_batch(X_real, y_real)
# generate 'fake' examples
X_fake, y_fake = generate_fake_samples(half_batch)
# update discriminator on fake samples
_, fake_acc = model.train_on_batch(X_fake, y_fake)
# summarize performance
print('>%d real=%.0f%% fake=%.0f%%' % (i+1, real_acc*100, fake_acc*100))

# define the discriminator model model
= define_discriminator()
# load image data

```

```
dataset = load_real_samples()  
# fit the model  
train_discriminator(model, dataset)
```

**Output:**

```
>1 real=26% fake=60%  
>2 real=38% fake=84%  
>3 real=35% fake=91%  
>4 real=34% fake=95%  
>5 real=36% fake=100%  
>6 real=41% fake=99%  
>7 real=31% fake=100%  
>8 real=51% fake=100%  
>9 real=45% fake=100%  
>10 real=42% fake=100%  
>11 real=48% fake=100%  
>12 real=47% fake=100%  
>13 real=52% fake=100%  
>14 real=57% fake=100%  
>15 real=64% fake=100%  
>16 real=66% fake=100%  
>17 real=59% fake=100%  
>18 real=70% fake=100%  
>19 real=80% fake=100%  
>20 real=77% fake=100%  
>21 real=83% fake=100%  
>22 real=81% fake=100%  
>23 real=82% fake=100%  
>24 real=92% fake=100%  
>25 real=91% fake=100%  
>26 real=91% fake=100%  
>27 real=91% fake=100%  
>28 real=95% fake=100%  
>29 real=95% fake=100%  
>30 real=95% fake=100%
```



```
>71 real=100% fake=100%
>72 real=100% fake=100%
>73 real=100% fake=100%
>74 real=100% fake=100%
>75 real=100% fake=100%
>76 real=100% fake=100%
>77 real=100% fake=100%
>78 real=100% fake=100%
>79 real=100% fake=100%
>80 real=100% fake=100%
>81 real=100% fake=100%
>82 real=100% fake=100%
>83 real=100% fake=100%
>84 real=100% fake=100%
>85 real=100% fake=100%
>86 real=100% fake=100%
>87 real=100% fake=100%
>88 real=100% fake=100%
>89 real=100% fake=100%
>90 real=100% fake=100%
>91 real=100% fake=100%
>92 real=100% fake=100%
>93 real=100% fake=100%
>94 real=100% fake=100%
>95 real=100% fake=100%
>96 real=100% fake=100%
>97 real=100% fake=100%
>98 real=100% fake=100%
>99 real=100% fake=100%
>100 real=100% fake=100%
```

## Practical 8

**Aim:** Write a program to implement a simple form of a recurrent neural network.

**1.(4-to-1 RNN) to show that the quantity of rain on a certain day also depends on the values of the previous day**

**Source Code:**

```
import numpy as np
import tensorflow as tf

# Define the training data
rainfall_data = np.array([[0.2, 0.3, 0.1, 0.5, 0.4],
                          [0.1, 0.4, 0.5, 0.2, 0.3],
                          [0.3, 0.2, 0.4, 0.3, 0.1],
                          [0.4, 0.1, 0.3, 0.4, 0.2],
                          [0.5, 0.5, 0.2, 0.1, 0.5]])

# Prepare the input and output data input_data = rainfall_data[:, :-1]
# Previous four days' rainfall values output_data = rainfall_data[:, -1]
# Current day's rainfall value

# Define the RNN model model
= tf.keras.Sequential([
    tf.keras.layers.SimpleRNN(10, input_shape=(4, 1)),
    tf.keras.layers.Dense(1)
])

# Compile the model
model.compile(loss='mse', optimizer='adam')

# Train the model
model.fit(np.expand_dims(input_data, axis=2), output_data, epochs=100, batch_size=1)

# Predict the rainfall for a new day new_input = np.array([[0.3, 0.2, 0.1, 0.4]]) #
Previous four days' rainfall values predicted_rainfall =
model.predict(np.expand_dims(new_input, axis=2)) print("Predicted rainfall for
the new day:", predicted_rainfall[0][0])
```

### Output:

```
5/5 [=====] - 0s 12ms/step - loss: 0.0024
Epoch 94/100
5/5 [=====] - 0s 10ms/step - loss: 0.0025
Epoch 95/100
5/5 [=====] - 0s 9ms/step - loss: 0.0024
Epoch 96/100
5/5 [=====] - 0s 15ms/step - loss: 0.0024
Epoch 97/100
5/5 [=====] - 0s 11ms/step - loss: 0.0024
Epoch 98/100
5/5 [=====] - 0s 10ms/step - loss: 0.0023
Epoch 99/100
5/5 [=====] - 0s 10ms/step - loss: 0.0023
Epoch 100/100
5/5 [=====] - 0s 8ms/step - loss: 0.0023
1/1 [=====] - 1s 520ms/step
Predicted rainfall for the new day: 0.33681476
```

### Source Code:

```
import pandas as pd
import numpy as np
import tensorflow as tf

from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt

data = pd.read_csv("/content/sample_data/training.txt", delimiter="\t", names=["label", "text"])

X_train, X_test, y_train, y_test = train_test_split(data["text"], data["label"], test_size=0.2,
                                                    random_state=42)

tokenizer = Tokenizer(num_words=5000, oov_token="<OOV>")
tokenizer.fit_on_texts(X_train)

X_train_seq = tokenizer.texts_to_sequences(X_train)
X_test_seq = tokenizer.texts_to_sequences(X_test)

max_length = 100
X_train_pad = pad_sequences(X_train_seq, maxlen=max_length, padding="post", truncating="post")
X_test_pad = pad_sequences(X_test_seq, maxlen=max_length, padding="post", truncating="post")

model = tf.keras.models.Sequential([tf.keras.layers.Embedding(input_dim=5000,
                                                                output_dim=32,
                                                                input_length=max_length),
                                    tf.keras.layers.LSTM(units=64, dropout=0.2,
                                                         recurrent_dropout=0.2),
                                    tf.keras.layers.Dense(1, activation="sigmoid")])
```

```

model.compile(optimizer="adam", loss="binary_crossentropy", metrics=["accuracy"])

history = model.fit(X_train_pad, y_train, epochs=10, batch_size=32, validation_split=0.1)

loss, accuracy = model.evaluate(X_test_pad, y_test)
print("Test loss:", loss) print("Test
accuracy:", accuracy)


plt.plot(history.history["accuracy"], label="Training accuracy")
plt.plot(history.history["val_accuracy"], label="Validation accuracy")
plt.xlabel("Epoch") plt.ylabel("Accuracy")
plt.legend()
plt.show()

predictions = model.predict(X_test_pad)

index = np.random.randint(0, len(X_test_pad)) text =
tokenizer.sequences_to_texts([X_test_pad[index]])[0] label
= y_test.values[index] prediction = predictions[index][0]
print("Text:", text) print("Actual
label:", label)
print("Predicted label:", round(prediction))

```

## Output:



## 2.LSTM for sentiment analysis on datasets like UMICH SI650 for similar.

### Source Code:

```
from __future__ import division, print_function
from keras.layers.core import Dense, Activation
from keras.layers import Embedding from
keras.layers import LSTM from keras.models import
Sequential from keras.preprocessing import
sequence from sklearn.model_selection import
train_test_split import collections import nltk import
numpy as np nltk.download('punkt')
```

```
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!
True
```

```
maxlen = 0
word_freqs = collections.Counter()
num_recs = 0

ftrain = open("/content/umich-sentiment-train.txt", "r") # Open the file in text mode

for line in ftrain:    label, sentence =
line.strip().split("\t")
    words = nltk.word_tokenize(sentence.lower())

    if len(words) > maxlen:
maxlen = len(words)

    for word in words:
        word_freqs[word] += 1

    num_recs += 1

ftrain.close()

# print some statistics about our data, that will drive our parameters print("maxlen:
%d, vocab size: %d" % (maxlen, len(word_freqs)))
```

```
maxlen: 42, vocab size: 2268
```

## Practical 9

**Aim:** Write a program for object detection from the image/video.

**Source Code:**

```
import torch import  
torchvision import  
pycocotools  
from PIL import Image
```

```
holiday = Image.open("/content/kids.jpg").convert('RGB')  
holiday
```



```
kids_playing = Image.open("/content/kids.jpg").convert('RGB') kids_playing
```





```
holiday_tensor_int = pil_to_tensor(holiday)
kids_playing_tensor_int = pil_to_tensor(kids_playing)
```

```
holiday_tensor_int.shape, kids_playing_tensor_int.shape
```

```
(torch.Size([3, 1120, 2016]), torch.Size([3, 1120, 2016]))
```

```
holiday_tensor_int = holiday_tensor_int.unsqueeze(dim=0)
kids_playing_tensor_int = kids_playing_tensor_int.unsqueeze(dim=0)
```

```
holiday_tensor_int.shape, kids_playing_tensor_int.shape
```

```
(torch.Size([1, 3, 1120, 2016]), torch.Size([1, 3, 1120, 2016]))
```

```
print(holiday_tensor_int.min(), holiday_tensor_int.max())
```

```
tensor(0, dtype=torch.uint8) tensor(255, dtype=torch.uint8)
```

```
holiday_tensor_float = holiday_tensor_int / 255.0
kids_playing_tensor_float = kids_playing_tensor_int / 255.0
```

```
print(holiday_tensor_float.min(), holiday_tensor_float.max())
```





```

[1177.9301, 492.4050, 1205.8940, 528.0626],
[ 795.9549, 438.3310, 912.4966, 558.4469],
[1182.7083, 497.9762, 1204.9606, 525.6568],
[ 929.5096, 366.9445, 1041.1884, 543.0920],
[1521.7596, 602.9040, 1597.0552, 630.3387],
[ 770.5167, 504.3984, 1049.3455, 681.6665],
[ 386.7892, 435.2683, 1089.5959, 668.9307],
[1782.7791, 704.2247, 1842.4552, 764.6519],
[ 350.5378, 437.2085, 483.5045, 680.3541],
[ 940.6858, 425.4261, 1053.9584, 565.6411],
[1192.0833, 479.0542, 1253.0463, 570.6131],
[1201.1323, 486.0803, 1250.7799, 566.1834],
[ 420.6662, 341.7614, 1061.7148, 632.9488],
[1203.5992, 484.4836, 1267.4409, 605.7438],
[ 640.4301, 203.6466, 753.3429, 671.9545]],
grad_fn=<StackBackward0>),
'labels': tensor([ 1, 1, 1, 1, 1, 37, 1, 37, 62, 62, 37, 39, 62, 15, 62, 62, 34, 62,
62, 37, 62, 15, 62, 1, 44, 62, 62, 62, 47, 62, 16, 62, 37, 62, 15, 15,
15, 53, 15, 1, 43, 39, 62, 43, 62]),
'scores': tensor([0.9998, 0.9998, 0.9998, 0.9997, 0.9993, 0.9628, 0.8721, 0.8368, 0.8039,
0.7754, 0.5970, 0.5537, 0.5425, 0.5032, 0.4730, 0.3431, 0.3246, 0.3053,
0.2679, 0.2108, 0.2053, 0.1895, 0.1618, 0.1466, 0.1383, 0.1331, 0.1305,
0.1283, 0.1064, 0.0995, 0.0917, 0.0902, 0.0869, 0.0832, 0.0803, 0.0801,
0.0789, 0.0658, 0.0654, 0.0615, 0.0596, 0.0567, 0.0559, 0.0543, 0.0504],
grad_fn=<IndexBackward0>))}]

```

```

holiday_preds[0]["boxes"] = holiday_preds[0]["boxes"][holiday_preds[0]["scores"] > 0.8]
holiday_preds[0]["labels"] = holiday_preds[0]["labels"][holiday_preds[0]["scores"] > 0.8]
holiday_preds[0]["scores"] = holiday_preds[0]["scores"][holiday_preds[0]["scores"] > 0.8]

```

holiday\_preds

```

[{'boxes': tensor([[1565.8942, 471.7659, 1811.0688, 927.4228],
[ 120.9253, 363.1345, 421.6961, 1105.6018],
[ 509.4385, 458.9574, 693.5939, 1001.9330],
[1181.9674, 477.7082, 1433.4752, 982.7266],
[ 741.7849, 462.7850, 981.6406, 1034.2094],
[1784.6957, 704.5729, 1842.0870, 763.6006],
[ 948.3671, 411.9908, 1111.9010, 652.2377],
[ 302.4559, 725.4138, 375.1870, 802.2778],
[ 372.9771, 474.3299, 481.1460, 667.6293]]),
grad_fn=<IndexBackward0>),
'labels': tensor([ 1, 1, 1, 1, 1, 37, 1, 37, 62]),
'scores': tensor([0.9998, 0.9998, 0.9998, 0.9997, 0.9993, 0.9628, 0.8721, 0.8368, 0.8039],
grad_fn=<IndexBackward0>))}]

```

```

kids_preds = object_detection_model(kids_playing_tensor_float)

```

kids\_preds

```
[{'boxes': tensor([[1565.8942, 471.7659, 1811.0688, 927.4228],
[ 120.9253, 363.1345, 421.6961, 1105.6018],
[ 509.4385, 458.9574, 693.5939, 1001.9330],
[1181.9674, 477.7082, 1433.4752, 982.7266],
[ 741.7849, 462.7850, 981.6406, 1034.2094],
[1784.6957, 704.5729, 1842.0870, 763.6006],
[ 948.3671, 411.9908, 1111.9010, 652.2377],
[ 302.4559, 725.4138, 375.1870, 802.2778],
[ 372.9771, 474.3299, 481.1460, 667.6293],
[ 397.3616, 331.0306, 480.6055, 666.0760],
[1729.4586, 588.4355, 1772.0046, 645.2954],
[1194.8232, 480.7062, 1270.5952, 596.5611],
[ 825.9108, 355.1837, 1008.1398, 690.1427],
[ 949.6096, 488.7204, 1180.6737, 735.1190],
[ 433.1131, 306.8869, 500.2051, 659.2569],
[ 798.1258, 433.2409, 903.2397, 500.1232],
[ 304.1548, 724.5444, 375.0820, 802.0087],
[ 893.0400, 368.0695, 1166.3256, 716.5153],
[ 938.7700, 502.0295, 1028.8362, 685.3348],
[1772.8767, 591.4392, 1787.5862, 621.2263],
[ 789.2773, 398.8733, 956.6437, 666.9479],
[ 861.0916, 359.2489, 1167.0411, 720.7684],
[ 838.9109, 345.3440, 1028.7206, 543.1509],
[ 972.9252, 432.1348, 1084.3429, 599.9109],
[1178.8203, 483.0518, 1207.9275, 526.8611],
[ 928.6594, 351.2741, 1027.0938, 445.9928],
[1148.1388, 248.4578, 1298.6729, 619.1220],
[ 806.9293, 383.5974, 927.6639, 583.2614],
[1184.6785, 483.6467, 1206.5623, 518.6398],
[ 928.0013, 402.2428, 1055.2094, 682.7582],
[1177.9301, 492.4050, 1205.8940, 528.0626],
[ 795.9549, 438.3310, 912.4966, 558.4469],
[1182.7083, 497.9762, 1204.9606, 525.6568],
```

```

[1102.7005, 437.5702, 1204.5000, 525.0500],
[ 929.5096, 366.9445, 1041.1884, 543.0920],
[1521.7596, 602.9040, 1597.0552, 630.3387],
[ 770.5167, 504.3984, 1049.3455, 681.6665],
[ 386.7892, 435.2683, 1089.5959, 668.9307],
[1782.7791, 704.2247, 1842.4552, 764.6519],
[ 350.5378, 437.2085, 483.5045, 680.3541],
[ 940.6858, 425.4261, 1053.9584, 565.6411],
[1192.0833, 479.0542, 1253.0463, 570.6131],
[1201.1323, 486.0803, 1250.7799, 566.1834],
[ 420.6662, 341.7614, 1061.7148, 632.9488],
[1203.5992, 484.4836, 1267.4409, 605.7438],
[ 640.4301, 203.6466, 753.3429, 671.9545]],
grad_fn=<StackBackward0>),
'labels': tensor([ 1,  1,  1,  1,  1, 37,  1, 37, 62, 62, 37, 39, 62, 15, 62, 62, 34, 62,
62, 37, 62, 15, 62,  1, 44, 62, 62, 62, 47, 62, 16, 62, 37, 62, 15, 15,
15, 53, 15,  1, 43, 39, 62, 43, 62]),
'scores': tensor([0.9998, 0.9998, 0.9998, 0.9997, 0.9993, 0.9628, 0.8721, 0.8368, 0.8039,
0.7754, 0.5970, 0.5537, 0.5425, 0.5032, 0.4730, 0.3431, 0.3246, 0.3053,
0.2679, 0.2108, 0.2053, 0.1895, 0.1618, 0.1466, 0.1383, 0.1331, 0.1305,
0.1283, 0.1064, 0.0995, 0.0917, 0.0902, 0.0869, 0.0832, 0.0803, 0.0801,
0.0789, 0.0658, 0.0654, 0.0615, 0.0596, 0.0567, 0.0559, 0.0543, 0.0504],
grad_fn=<IndexBackward0>))}]

```

```

kids_preds[0]["boxes"] = kids_preds[0]["boxes"][kids_preds[0]["scores"] > 0.8] kids_preds[0]["labels"]
= kids_preds[0]["labels"][kids_preds[0]["scores"] > 0.8] kids_preds[0]["scores"] =
kids_preds[0]["scores"][kids_preds[0]["scores"] > 0.8]

```

kids\_preds

```

[{'boxes': tensor([[1565.8942, 471.7659, 1811.0688, 927.4228],
[ 120.9253, 363.1345, 421.6961, 1105.6018],
[ 509.4385, 458.9574, 693.5939, 1001.9330],
[1181.9674, 477.7082, 1433.4752, 982.7266],
[ 741.7849, 462.7850, 981.6406, 1034.2094],
[1784.6957, 704.5729, 1842.0870, 763.6006],
[ 948.3671, 411.9908, 1111.9010, 652.2377],
[ 302.4559, 725.4138, 375.1870, 802.2778],
[ 372.9771, 474.3299, 481.1460, 667.6293]],
grad_fn=<IndexBackward0>),
'labels': tensor([ 1,  1,  1,  1,  1, 37,  1, 37, 62]),
'scores': tensor([0.9998, 0.9998, 0.9998, 0.9997, 0.9993, 0.9628, 0.8721, 0.8368, 0.8039],
grad_fn=<IndexBackward0>)}]

```

```
from pycocotools.coco import COCO
```

```
annFile='/content/instances_val2017.json'
```

```
coco=COCO(annFile)
```

```
loading annotations into memory...
Done (t=0.00s)
creating index...
index created!
```

```
holiday_labels = coco.loadCats(holiday_preds[0]["labels"].numpy())
```

```
holiday_labels
```

```
[{'supercategory': 'person', 'id': 1, 'name': 'person'},
 {'supercategory': 'person', 'id': 1, 'name': 'person'},
 {'supercategory': 'person', 'id': 1, 'name': 'person'},
 {'supercategory': 'person', 'id': 1, 'name': 'person'},
 {'supercategory': 'person', 'id': 1, 'name': 'person'},
 {'supercategory': 'sports', 'id': 37, 'name': 'sports ball'},
 {'supercategory': 'person', 'id': 1, 'name': 'person'},
 {'supercategory': 'sports', 'id': 37, 'name': 'sports ball'},
 {'supercategory': 'furniture', 'id': 62, 'name': 'chair'}]
```

```
kids_labels = coco.loadCats(kids_preds[0]["labels"].numpy())
```

```
kids_labels
```

```
[{'supercategory': 'person', 'id': 1, 'name': 'person'},
 {'supercategory': 'person', 'id': 1, 'name': 'person'},
 {'supercategory': 'person', 'id': 1, 'name': 'person'},
 {'supercategory': 'person', 'id': 1, 'name': 'person'},
 {'supercategory': 'person', 'id': 1, 'name': 'person'},
 {'supercategory': 'sports', 'id': 37, 'name': 'sports ball'},
 {'supercategory': 'person', 'id': 1, 'name': 'person'},
 {'supercategory': 'sports', 'id': 37, 'name': 'sports ball'},
 {'supercategory': 'furniture', 'id': 62, 'name': 'chair'}]
```

```
from torchvision.utils import draw_bounding_boxes
```

```
holiday_annot_labels = [{"{}-{:2f}".format(label["name"], prob) for label, prob in zip(holiday_labels,
holiday_preds[0]["scores"].detach().numpy())}]
```

```

holiday_output = draw_bounding_boxes(image=holiday_tensor_int[0],
boxes=holiday_preds[0]["boxes"],          labels=holiday_annot_labels,
                                     colors=["red" if label["name"]=="person" else "green" for label in holiday_labels],
                                     width=2
)

```

```

holiday_output.shape

```

```

torch.Size([3, 1120, 2016])

```

```


from torchvision.transforms.functional import to_pil_image

```

```

to_pil_image(holiday_output)

```



```

from torchvision.utils import draw_bounding_boxes

```

```

kids_annot_labels = [{"{}-{:2f}".format(label["name"], prob) for label, prob in zip(kids_labels,
kids_preds[0]["scores"].detach().numpy())]

```

```

kids_output = draw_bounding_boxes(image=kids_playing_tensor_int[0],
boxes=kids_preds[0]["boxes"],          labels=kids_annot_labels,
                                     colors=["red" if label["name"]=="person" else "green" for label in kids_labels],

```

```
width=2,  
font_size=16,  
fill=True  
)
```

```
to_pil_image(kids_output)
```

